

```

In [ ]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder, StandardScaler, MinMaxScaler
from sklearn.impute import SimpleImputer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.cluster import KMeans
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import mean_squared_error, r2_score, confusion_matrix, classification_report

file_path = 'LABTest - LABTest.csv'

# Loading the data
data = pd.read_csv(file_path)

# Display the first few rows of the dataset
data.head()

```

```

Out[ ]:

```

	Administrative	Administrative_Duration	Informational	Informational_Duration	ProductRelated
0	0	0.0	0	0.0	1
1	0	0.0	0	0.0	2
2	0	0.0	0	0.0	1
3	0	0.0	0	0.0	2
4	0	0.0	0	0.0	10

```

In [ ]: # Exploratory analysis

# Display basic information about the dataset
data.info()

# Display statistical summary of the dataset
data.describe()

```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12330 entries, 0 to 12329
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Administrative                        12330 non-null  int64
1   Administrative_Duration               12330 non-null  float64
2   Informational                        12330 non-null  int64
3   Informational_Duration               12330 non-null  float64
4   ProductRelated                      12330 non-null  int64
5   ProductRelated_Duration             12330 non-null  float64
6   BounceRates                         12330 non-null  float64
7   ExitRates                          12330 non-null  float64
8   PageValues                         12330 non-null  float64
9   SpecialDay                         12330 non-null  float64
10  Month                              12330 non-null  object
11  OperatingSystems                   12330 non-null  int64
12  Browser                           12330 non-null  int64
13  Region                            12330 non-null  int64
14  TrafficType                       12330 non-null  int64
15  VisitorType                       12330 non-null  object
16  Weekend                           12330 non-null  bool
17  Revenue                           12330 non-null  bool
dtypes: bool(2), float64(7), int64(7), object(2)
memory usage: 1.5+ MB
```

Out []:

	Administrative	Administrative_Duration	Informational	Informational_Duration	ProductRelat
count	12330.000000	12330.000000	12330.000000	12330.000000	12330.0000
mean	2.315166	80.818611	0.503569	34.472398	31.7314
std	3.321784	176.779107	1.270156	140.749294	44.4755
min	0.000000	0.000000	0.000000	0.000000	0.0000
25%	0.000000	0.000000	0.000000	0.000000	7.0000
50%	1.000000	7.500000	0.000000	0.000000	18.0000
75%	4.000000	93.256250	0.000000	0.000000	38.0000
max	27.000000	3398.750000	24.000000	2549.375000	705.0000

```
In [ ]: # Check for missing values
missing_values = data.isnull().sum()
print(missing_values[missing_values > 0])

# Fill missing values with the mean for numerical columns and mode for categorical columns
for column in data.columns:
    if data[column].dtype == 'object':
        data[column].fillna(data[column].mode()[0], inplace=True)
    else:
        data[column].fillna(data[column].mean(), inplace=True)

Series([], dtype: int64)
```

```
In [ ]: # convert categorical values to numerical using onehotencoding using pandas getdummies
# Encode categorical variables using one-hot encoding
data = pd.get_dummies(data, drop_first=True)
```

```
In [ ]: # Preprocess the dataset using MinMaxScaler Normalization
from sklearn.preprocessing import MinMaxScaler
minmax = MinMaxScaler()
minmax.fit_transform(data)
data.head()
```

```
Out[ ]:      Administrative  Administrative_Duration  Informational  Informational_Duration  ProductRelated
0                0                0.0                0                0.0                1
1                0                0.0                0                0.0                2
2                0                0.0                0                0.0                1
3                0                0.0                0                0.0                2
4                0                0.0                0                0.0                10
```

5 rows × 27 columns



```
In [ ]: # split the dataset into train and split data
# The dataset consists of 18 attributes. The Revenue attribute can be used as the class label
# Administrative, Administrative Duration, Informational, Informational Duration, Product
# Related and Product Related Duration represent the number of different types of pages visited
# by the visitor in that session and total time spent in each of these page categories.
from sklearn.model_selection import train_test_split
X = data[['Administrative', 'Informational', 'ProductRelated', 'ProductRelated_Duration']]
Y = data['Revenue']
x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=42)
```

```
In [ ]: # Perform Classification using KNN classifier for atleast 4 different values of K

k_values = [2,3,4,5]

for i in k_values:
    modelknn = KNeighborsClassifier(n_neighbors=i, metric="euclidean")
    modelknn.fit(x_train, y_train)

# Predictions for KNN
y_train_pred_knn = modelknn.predict(x_train)
y_test_pred_knn = modelknn.predict(x_test)
```

```
In [ ]: # Perform Classification using NaiveBayes classifier using all possible parameters

# Build Naive Bayes classifier
nb = GaussianNB()
nb.fit(x_train, y_train)
```

```
# Predictions for Naive Bayes
y_train_pred_nb = nb.predict(x_train)
y_test_pred_nb = nb.predict(x_test)
```


```
In [ ]: # Cluster the visitors into 8 clusters ,visualize it and compute the silhouette score score V
inertias = []
# Performing Elbow Method to find the Best number of clusters
for i in range(1, 11):
    kmeans = KMeans(n_clusters=i, random_state=42)
    kmeans.fit(data)
    inertias.append(kmeans.inertia_)

# plotting the possible k values
plt.figure(figsize=(8, 4))
plt.plot(range(1, 11), inertias, marker='o')
plt.title('Elbow Method')
plt.xlabel('Number of Clusters')
plt.ylabel('Inertia')
plt.grid(True)
plt.show()

kmeans = KMeans(n_clusters=4, random_state=42)
kmeans.fit(data)

labels = kmeans.labels_

data['Cluster'] = labels
```



By interpreting the Elbow graph we could understand that the elbow is formed at 2 but it would cause many errors at that datapoint hence we are choosing 3 as the value for the numbers

```

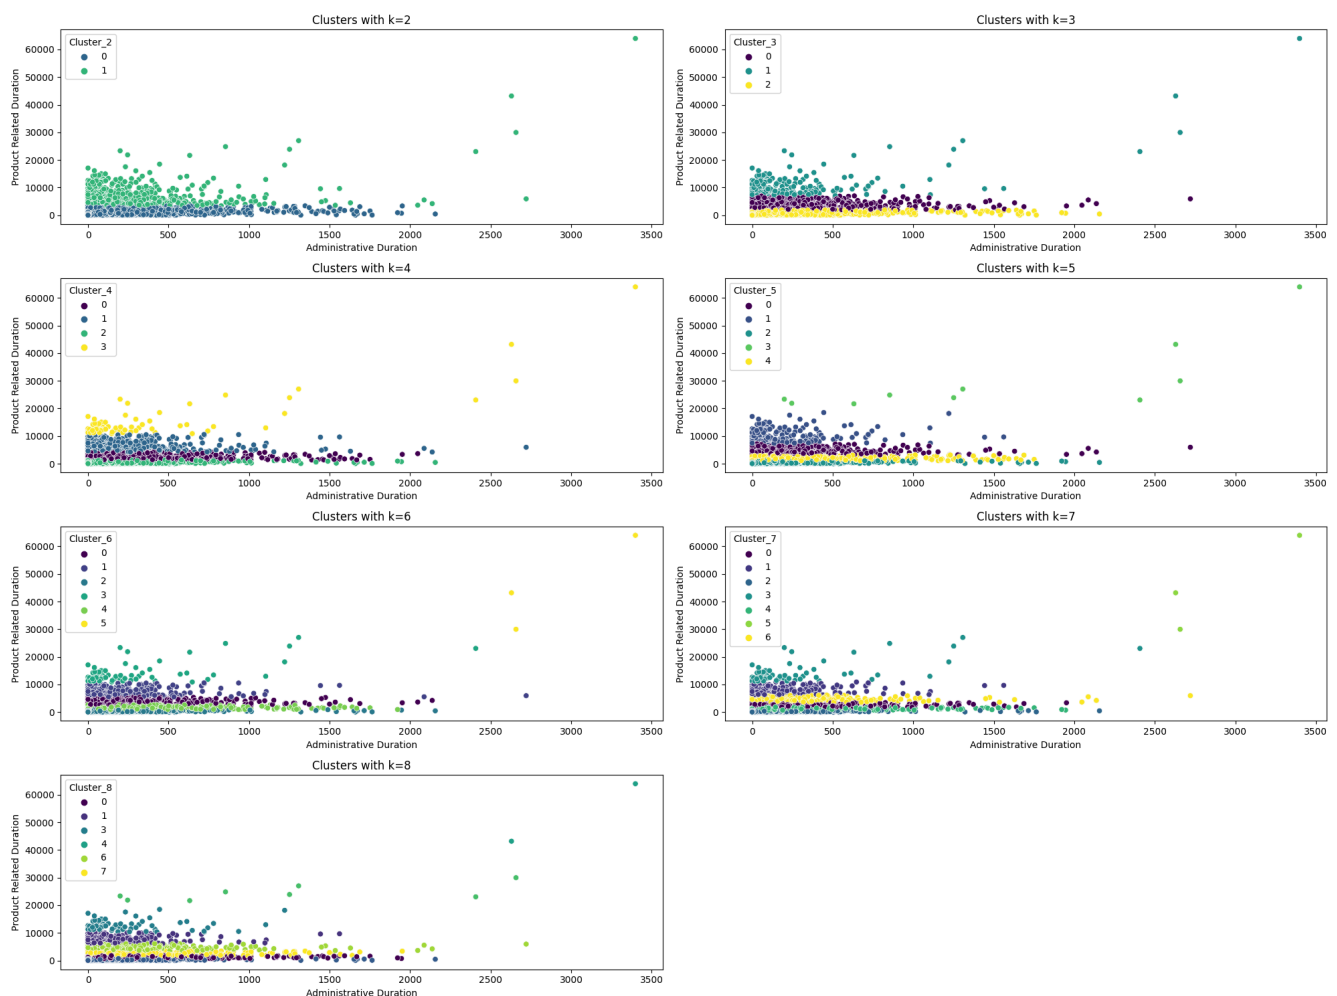
In [ ]: # providing all possible k values
kcluster_values = [2, 3, 4, 5, 6, 7, 8]
plt.figure(figsize=(20, 15))

# plotting the clusters for the 7 values of k
for idx, k in enumerate(kcluster_values):
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(data)
    labels = kmeans.labels_
    data[f'Cluster_{k}'] = labels

    plt.subplot(4, 2, idx + 1)
    sns.scatterplot(data=data, x='Administrative_Duration', y='ProductRelated_Duration', hue=f'Cluster_{k}')
    plt.title(f'Clusters with k={k}')
    plt.xlabel('Administrative Duration')
    plt.ylabel('Product Related Duration')

plt.tight_layout()
plt.show()

```



```

In [ ]: # According to the question we are taking into account for validating wheather 8 clusters are valid
# We are taking into account for the following metrics
# Silhouette score
# Calculating silhouettes score for K= 3 clustering
kmeans = KMeans(n_clusters=k, random_state=42)
kmeans.fit(data)

```

```

cluster_3_score = silhouette_score(X,data['Cluster_3'])
print("Silhouettes Score for K = 3",cluster_3_score)

# calculating silhouettes score for K= 8 clustering
kmeans = KMeans(n_clusters=k, random_state=42)
kmeans.fit(data)
cluster_3_score = silhouette_score(X,data['Cluster_8'])
print("Silhouettes Score for K = 8",cluster_3_score)

```

0.7108026008602492

0.6132905480245872

```

In [ ]: # Display Classification report and confusion matrix for both classifiers
# Evaluate KNN classifier
print('\nKNN Classifier:')
print(f'Accuracy on training set: {accuracy_score(y_train, y_train_pred_knn)}')
print(f'Accuracy on testing set: {accuracy_score(y_test, y_test_pred_knn)}')
print('Classification Report (Testing set):')
print(classification_report(y_test, y_test_pred_knn))
print('Confusion Matrix (Testing set):')
print(confusion_matrix(y_test, y_test_pred_knn))

print("\n-----")

# Evaluate Naive Bayes classifier
print('\nNaive Bayes Classifier:')
print(f'Accuracy on training set: {accuracy_score(y_train, y_train_pred_nb)}')
print(f'Accuracy on testing set: {accuracy_score(y_test, y_test_pred_nb)}')
print('Classification Report (Testing set):')
print(classification_report(y_test, y_test_pred_nb))
print('Confusion Matrix (Testing set):')
print(confusion_matrix(y_test, y_test_pred_nb))

```

KNN Classifier:

Accuracy on training set: 0.8577656123276561

Accuracy on testing set: 0.813463098134631

Classification Report (Testing set):

	precision	recall	f1-score	support
False	0.84	0.96	0.90	2055
True	0.25	0.06	0.10	411
accuracy			0.81	2466
macro avg	0.54	0.51	0.50	2466
weighted avg	0.74	0.81	0.76	2466

Confusion Matrix (Testing set):

```
[[1981  74]
 [ 386  25]]
```

Naive Bayes Classifier:

Accuracy on training set: 0.811536901865369

Accuracy on testing set: 0.810624493106245

Classification Report (Testing set):

	precision	recall	f1-score	support
False	0.85	0.94	0.89	2055
True	0.36	0.18	0.24	411
accuracy			0.81	2466
macro avg	0.61	0.56	0.57	2466
weighted avg	0.77	0.81	0.78	2466

Confusion Matrix (Testing set):

```
[[1925 130]
 [ 337  74]]
```

In []: *# Provide suitable inferences for the models created*

Inferences

Classification inference

For Naive bayes Accuracy on training set: 0.811536901865369 Accuracy on testing set: 0.810624493106245

For KNN Accuracy on training set: 0.8577656123276561 Accuracy on testing set: 0.813463098134631

knn has better accuracy for a very small variation precision is better for Naive bayes for true classification recall is better for KNN for true classification

Clustering Inference

We could see that the $k = 3$ has a silhouette score value than $k = 8$ hence $k = 3$ is a more accurate value for the number of clusters